

NLP Programming Tutorial 2 - Bigram Language Models

Graham Neubig
Nara Institute of Science and Technology (NAIST)

Review:

Calculating Sentence Probabilities

- We want the probability of

W = speech recognition system

- Represent this mathematically as:

$$P(|W| = 3, w_1 = \text{"speech"}, w_2 = \text{"recognition"}, w_3 = \text{"system"}) =$$

$$P(w_1 = \text{"speech"} \mid w_0 = \text{"<s>"})$$

$$* P(w_2 = \text{"recognition"} \mid w_0 = \text{"<s>"}, w_1 = \text{"speech"})$$

$$* P(w_3 = \text{"system"} \mid w_0 = \text{"<s>"}, w_1 = \text{"speech"}, w_2 = \text{"recognition"})$$

$$* P(w_4 = \text{"</s>"} \mid w_0 = \text{"<s>"}, w_1 = \text{"speech"}, w_2 = \text{"recognition"}, w_3 = \text{"system"})$$

NOTE:

sentence start <s> and end </s> symbol

NOTE:

$$P(w_0 = \text{"<s>"}) = 1$$

Incremental Computation

- Previous equation can be written:

$$P(W) = \prod_{i=1}^{|W|+1} P(w_i | w_0 \dots w_{i-1})$$

- **Unigram model** ignored context:

$$P(w_i | w_0 \dots w_{i-1}) \approx P(w_i)$$

Unigram Models Ignore Word Order!

- Ignoring context, probabilities are the same:

$$P_{\text{uni}}(w=\text{speech recognition system}) = \\ P(w=\text{speech}) * P(w=\text{recognition}) * P(w=\text{system}) * P(w=\text{</s>})$$

=

$$P_{\text{uni}}(w=\text{system recognition speech}) = \\ P(w=\text{speech}) * P(w=\text{recognition}) * P(w=\text{system}) * P(w=\text{</s>})$$

Unigram Models Ignore Agreement!

- Good sentences (words agree):

$$P_{\text{uni}}(w=\text{i am}) = P(w=\text{i}) * P(w=\text{am}) * P(w=</s>)$$

$$P_{\text{uni}}(w=\text{we are}) = P(w=\text{we}) * P(w=\text{are}) * P(w=</s>)$$

- Bad sentences (words don't agree)

$$P_{\text{uni}}(w=\text{we am}) = P(w=\text{we}) * P(w=\text{am}) * P(w=</s>)$$

$$P_{\text{uni}}(w=\text{i are}) = P(w=\text{i}) * P(w=\text{are}) * P(w=</s>)$$

But no penalty because probabilities are independent!

Solution: Add More Context!

- **Unigram** model ignored context:

$$P(w_i | w_0 \dots w_{i-1}) \approx P(w_i)$$

- **Bigram** model adds one word of context

$$P(w_i | w_0 \dots w_{i-1}) \approx P(w_i | w_{i-1})$$

- **Trigram** model adds two words of context

$$P(w_i | w_0 \dots w_{i-1}) \approx P(w_i | w_{i-2} w_{i-1})$$

- Four-gram, five-gram, six-gram, etc...

Maximum Likelihood Estimation of n-gram Probabilities

- Calculate counts of **n word** and **n-1 word** strings

$$P(w_i | w_{i-n+1} \dots w_{i-1}) = \frac{c(w_{i-n+1} \dots w_i)}{c(w_{i-n+1} \dots w_{i-1})}$$

i live in **osaka** . </s>

i am a graduate student . </s>

my school is in **nara** . </s>

$n=2 \rightarrow$

$$P(\text{osaka} | \text{in}) = c(\text{in osaka}) / c(\text{in}) = 1 / 2 = 0.5$$
$$P(\text{nara} | \text{in}) = c(\text{in nara}) / c(\text{in}) = 1 / 2 = 0.5$$

Still Problems of Sparsity

- When n-gram frequency is 0, probability is 0

$$P(\text{osaka} \mid \text{in}) = c(\text{i osaka})/c(\text{in}) = 1 / 2 = 0.5$$

$$P(\text{nara} \mid \text{in}) = c(\text{i nara})/c(\text{in}) = 1 / 2 = 0.5$$

$$P(\text{school} \mid \text{in}) = c(\text{in school})/c(\text{in}) = 0 / 2 = \mathbf{0!!}$$

- Like unigram model, we can use linear interpolation

Bigram: $P(w_i \mid w_{i-1}) = \lambda_2 P_{ML}(w_i \mid w_{i-1}) + (1 - \lambda_2) P(w_i)$

Unigram: $P(w_i) = \lambda_1 P_{ML}(w_i) + (1 - \lambda_1) \frac{1}{N}$

Choosing Values of λ : Grid Search

- One method to choose λ_2, λ_1 : try many values

$$\lambda_2 = 0.95, \lambda_1 = 0.95$$

$$\lambda_2 = 0.95, \lambda_1 = 0.90$$

$$\lambda_2 = 0.95, \lambda_1 = 0.85$$

...

$$\lambda_2 = 0.95, \lambda_1 = 0.05$$

$$\lambda_2 = 0.90, \lambda_1 = 0.95$$

$$\lambda_2 = 0.90, \lambda_1 = 0.90$$

...

$$\lambda_2 = 0.05, \lambda_1 = 0.10$$

$$\lambda_2 = 0.05, \lambda_1 = 0.05$$

Problems:

Too many options

→ Choosing takes time!

Using same λ for all n-grams

→ There is a smarter way!

Context Dependent Smoothing

High frequency word: “Tokyo”

$c(\text{Tokyo city}) = 40$
 $c(\text{Tokyo is}) = 35$
 $c(\text{Tokyo was}) = 24$
 $c(\text{Tokyo tower}) = 15$
 $c(\text{Tokyo port}) = 10$

...

Most 2-grams already exist
→ Large λ is better!

Low frequency word: “Tottori”

$c(\text{Tottori is}) = 2$
 $c(\text{Tottori city}) = 1$
 $c(\text{Tottori was}) = 0$

Many 2-grams will be missing
→ Small λ is better!

- Make the interpolation depend on the context

$$P(w_i | w_{i-1}) = \lambda_{w_{i-1}} P_{ML}(w_i | w_{i-1}) + (1 - \lambda_{w_{i-1}}) P(w_i)$$

Witten-Bell Smoothing

- One of the many ways to choose $\lambda_{w_{i-1}}$

$$\lambda_{w_{i-1}} = 1 - \frac{u(w_{i-1})}{u(w_{i-1}) + c(w_{i-1})}$$

$u(w_{i-1})$ = number of unique words after w_{i-1}

- For example:

$$\begin{array}{ll} c(\text{Tottori is}) = 2 & c(\text{Tottori city}) = 1 \\ c(\text{Tottori}) = 3 & u(\text{Tottori}) = 2 \end{array}$$

$$\lambda_{\text{Tottori}} = 1 - \frac{2}{2 + 3} = 0.6$$

$$\begin{array}{ll} c(\text{Tokyo city}) = 40 & c(\text{Tokyo is}) = 35 \dots \\ c(\text{Tokyo}) = 270 & u(\text{Tokyo}) = 30 \end{array}$$

$$\lambda_{\text{Tokyo}} = 1 - \frac{30}{30 + 270} = 0.9$$

Programming Techniques

Inserting into Arrays

- To calculate n-grams easily, you may want to:

```
my_words = ["this", "is", "a", "pen"]
```



```
my_words = ["<s>", "this", "is", "a", "pen", "</s>"]
```

- This can be done with:

```
my_words.append("</s>") # Add to the end
```

```
my_words.insert(0, "<s>") # Add to the beginning
```

Removing from Arrays

- Given an n-gram with $w_{i-n+1} \dots w_i$, we may want the context $w_{i-n+1} \dots w_{i-1}$
- This can be done with:

```
my_ngram = "tokyo tower"  
my_words = my_ngram.split(" ") # Change into ["tokyo", "tower"]  
my_words.pop()                 # Remove the last element ("tower")  
my_context = " ".join(my_words) # Join the array back together  
print my_context
```

Exercise

Exercise

- Write two programs
 - train-bigram: Creates a bigram model
 - test-bigram: Reads a bigram model and calculates entropy on the test set
- Test train-bigram on test/02-train-input.txt
- Train the model on data/wiki-en-train.word
- Calculate entropy on data/wiki-en-test.word (if linear interpolation, test different values of λ_2)
- Challenge:
 - Use Witten-Bell smoothing (Linear interpolation is easier)
 - Create a program that works with *any* n (not just bi-gram)

train-bigram (Linear Interpolation)

create **map** *counts*, *context_counts*

for each *line* **in** the *training_file*

split *line* into an array of *words*

append “</s>” to the end and “<s>” to the beginning of *words*

for each *i* **in** 1 **to** length(*words*)-1 # Note: starting at 1, **after** <s>

counts[" $w_{i-1} w_i$ "] += 1 # Add bigram and bigram context

context_counts[" w_{i-1} "] += 1

counts[" w_i "] += 1 # Add unigram and unigram context

context_counts[""] += 1

open the *model_file* for writing

for each *ngram*, *count* **in** *counts*

split *ngram* into an array of *words* # " $w_{i-1} w_i$ " → {" w_{i-1} ", " w_i "}

remove the last element of *words* # {" w_{i-1} ", " w_i "} → {" w_{i-1} "}

join *words* into *context* # {" w_{i-1} "} → " w_{i-1} "

probability = *counts*[*ngram*]/*context_counts*[*context*]

print *ngram*, *probability* **to** *model_file*

test-bigram (Linear Interpolation)

$\lambda_1 = ???, \lambda_2 = ???, V = 1000000, W = 0, H = 0$

load model into *probs*

for each *line* **in** *test_file*

split *line* into an array of *words*

append “</s>” to the end and “<s>” to the beginning of *words*

for each *i* **in** 1 **to** length(*words*)-1 # Note: starting at 1, **after** <s>

$P1 = \lambda_1 probs[“w_i”] + (1 - \lambda_1) / V$ # Smoothed unigram probability

$P2 = \lambda_2 probs[“w_{i-1} w_i”] + (1 - \lambda_2) * P1$ # Smoothed bigram probability

$H += -\log_2(P2)$

$W += 1$

print “entropy = ”+ H/W

Thank You!